

ENCM 369 Winter 2023 Lab 6 for the Week of February 27

Steve Norman
Department of Electrical & Software Engineering
University of Calgary

February 2023

Administrative details

You may work with a partner on this assignment

If you choose to work with a partner, please make sure that both partners fully understand all parts of your assignment submission, and please follow the instructions regarding submission of your PDF document.

Partners must be in the same lab section. The reason for this rule to keep teaching assistant workloads balanced and to make it as easy as possible for teaching assistants to record marks.

Due Dates

The Due Date for this assignment is 6:00pm Friday, March 3.

The Late Due Date is 6:00pm Monday, March 6.

The penalty for handing in an assignment after the Due Date but before the Late Due Date is 3 marks. In other words, X/Y becomes $(X-3)/Y$ if the assignment is late. There will be no credit for assignments turned in after the Late Due Date; they will not be marked.

How to package and hand in your assignments

You must submit your work as a *single PDF file* to the D2L dropbox that will be set up for this assignment. The dropbox will be configured to accept only file per student, but you may replace that file as many times as you like before the due date.

See the Lab 1 instructions for more information about preparing and uploading your PDF file.

Important update for those working with a partner: Please submit only one PDF file for both students. On the cover page, put lab section, name and ICID information in this format:

Group Submission for [lab section]

Submitted by: [submitter's name]

UCID: [submitter's UCID]

Partner: [partner's name]

UCID: [partner's UCID]

Marking scheme

A	2 marks
B	4 marks
E	4 marks
F	4 marks
TOTAL	14 marks

Getting files for this lab

The files you need for this week's exercises are in a folder called `encm369w23lab06`, which has been uploaded to D2L in `.zip` format. Before starting work, you must download the `.zip` file, and extract its contents in a suitable place in the file system.

Exercise A: 12-bit two's complement and RISC-V instructions

Here is a RISC-V assembly language instruction to allocate 28 words on the stack:

```
addi    sp, sp, -112
```

Answer the following questions:

1. What is the machine code for the instruction? Show how you obtained your answer, and express your answer in base two.
2. Suppose before the instruction is run, `sp = 0x7fff_ed80`. A 32-bit adder will compute the new value for `sp`. What are the two 32-bit inputs to the adder, and what is the 32-bit output of the adder? Show how you obtained your answers, and express all your answers in base two.

What to Include in Your PDF Submission

Include answers to questions 1 and 2.

Exercise B: Integer addition examples

Read This First

This is a pencil and paper exercise. Don't use a computer or a calculator.

In this exercise, you get to add bit patterns as if you were an 8-bit processor, millions of times slower than electronic hardware.

What to Do

For each of the following pairs of hexadecimal patterns, write 8-bit binary patterns, and find the sum that would be generated by adding the patterns with an 8-bit adder.

PART	<i>a</i>	<i>b</i>	PART	<i>a</i>	<i>b</i>
I	0xb4	0xb3	III	0x78	0x0b
II	0xd0	0xe0	IV	0x35	0x2d

For each part, do the following after calculating the sum:

- First interpret the numbers as two's-complement signed values. Decide if the addition resulted in *signed* overflow.

- Then interpret the numbers as unsigned values. Decide if the addition resulted in *unsigned* overflow.

Show the work you did to generate the addition results and explain in each case how you decided whether overflow occurred.

What to Include in Your PDF Submission

Include answers and supporting work for all the additions.

Exercise C: Scary facts about real-world C integer types

Read This First

A key fact, emphasized in ENCM 369 lectures: Two's complement is by far the most common system for representing signed integers in computers. (This fact is not scary, but the next two are.)

Another key fact, also emphasized in ENCM 369 lectures: Typical implementations of C and C++ will allow integer arithmetic to generate results that do not make sense according to normal, everyday mathematics; these implementations will not do anything special to alert programmers or users that results probably do not make sense. For example, it is quite easy to get a C program to add two `int` variables with positive values and generate a sum that is apparently negative.

Here is a third key fact: C and C++ programming language standards *do not exactly specify* the numbers of bits that should be used for numerical types!

Back in the early days of C, there were three signed integer types: `short int`, `int`, and `long int`. (There were also various unsigned types, and `char`, which was usually an 8-bit type, but which was signed on some systems and unsigned on others.) It may seem weird, but for a given platform, typically one of the following was true:

- `short int` and `int` were the same size, while `long int` was wider than `int`;
- `short int` was narrower than `int`, while `long int` and `int` were the same size.

Later a type called `long long int` became available; this type was usually 64 bits wide. Figure 1 gives some size information for some C types on some important platforms.

Now consider software development in the early 1980's. It would not have been uncommon for a programmer to write C code for an expensive Unix workstation, get an executable built and tested, then later copy the C source code to a cheaper PC running MS-DOS and try to build an executable for the PC. If values of `ints` could be expected to go beyond the range of a 16-bit type, the PC version of the program could fail in ways that were never seen on Unix, and that were either mysterious or just annoying, depending on how much the programmer knew about how `int` values were represented on the two different platforms.

In 2023, a similar problem can occur if programmers move code from one platform to another and are not careful to check on the width of the `long int` type—the width is 64 bits for many current platforms, but only 32 bits for others.

(By the way, the designers of the Java programming language had experience with the problems of a given C or C++ type having different sizes on different platforms. For Java there are strict rules such as, “an `int` must be represented using 32-bit two's-complement,” and “a `long` must be represented using 64-bit two's-complement.”)

Figure 1: Sizes for signed integer types on various platforms.

Platform	short int	int	long int	long long int
Typical Unix minicomputer or workstation, starting in early 1980's	16 bits	32 bits	32 bits	n/a
PC with MS-DOS and typical C compiler, starting in early 1980's	16 bits	16 bits	32 bits	n/a
Linux system with x86 architecture	16 bits	32 bits	32 bits	64 bits
Linux system with x86-64 architecture	16 bits	32 bits	64 bits	64 bits
Current and recent macOS (64-bit)	16 bits	32 bits	64 bits	64 bits
Current and recent Microsoft Windows (both 32- and 64-bit), with Microsoft compilers	16 bits	32 bits	32 bits	64 bits
Cygwin64 running on top of Microsoft Windows	16 bits	32 bits	64 bits	64 bits

What to Do

Find the file `long_ints.c` in directory `encm369w23lab06/exC`

Do not yet make an executable and run it. Instead, carefully read the code in `long_ints.c` and develop predictions for the program output for the following cases: first, a platform on which `long int` is a 32-bit type; second, a platform on which `long int` is a 64-bit type. Here are a few facts which may be helpful:

$$\begin{aligned}
 2^{30} &= 1,073,741,824 \\
 2^{30} + 2^{28} &= 1,342,177,280 \\
 2^{31} + 2^{28} &= 2,415,919,104 \\
 2^{31} - 2^{28} &= 1,879,048,192
 \end{aligned}$$

You can test your prediction for 64-bit `long int` by building and running an executable on Cygwin64, 64-bit macOS, or a 64-bit Linux system.

For 32-bits, you can write a small RARS program that uses `add` to add the integers 1342177280 and 1073741824, then prints the sum.

What to Include in Your PDF Submission

There is nothing to include. If it's not clear to you why the two programs produce the output they do, you can check an explanation that will be posted on the course "lab solutions" web page on or before February 28.

Exercise D: 64-bit addition with 32-bit registers

Read This First

The material in Exercise C may raise a number of questions. One of them might be: How is a 64-bit integer type possible on a machine where all the general-purpose registers (GPRs) are 32 bits wide?

A partial answer is: If 64-bit integer variables are to be in GPRs, each of those variables will need two GPRs; an operation such as addition cannot be done with a single instruction, so will need a sequence of instructions instead.

Consider this situation on RISC-V: *i*, *j*, and *k* are all 64-bit signed integer variables in a C program. Suppose registers are allocated this way:

i: bits 31–0 in *s0*, bits 63–32 in *s1*
j: bits 31–0 in *s2*, bits 63–32 in *s3*
k: bits 31–0 in *s4*, bits 63–32 in *s5*

How could the C statement *k* = *i* + *j*; be implemented in assembly language?

Let's pretend that two untrue things are actually true. First, let's suppose that the RISC-V *add* instruction places the carry out of the most significant bit of the adder hardware into a special location within the processor called the *carry bit*. Second, let's suppose there is an instruction called *addc* ("add-with-carry") that is exactly like *add*, except that it takes the 0 or 1 in the carry bit as a carry into the least significant bit of the adder hardware. With those two pretended conditions in place, the 64-bit addition is simple:

```
add    s4, s0, s2 # get bits 31-0 of result
addc   s5, s1, s3 # get bits 63-32 of result with
                # appropriate carry from bit 31 to 32
```

So why did I bother to explain how to solve the problem using fictional properties of the RISC-V architecture? I did it because the carry bit and add-with-carry instructions are available on a wide range of important architectures, including x86-84 and ARM, so this kind of solution is useful to know about.

A similar solution is usually available to allow 32-bit integer addition on architectures with 16-bit GPRs.

What to Do

Consider the problem of implementing the 64-bit addition of the Read This First section with real RISC-V instructions.

1. Find a solution that is a sequence of four instructions—*add*, another *add*, *addi*, *bgeu*, *not* in that exact order. Hint: Think about the condition to detect unsigned overflow in addition—that happens if and only if the carry out of the most significant bit is 1.
2. Find another solution that is also four instructions in length, but without any branches. Hint: Think about what possible bit patterns *sltu* could put into its destination register.

Note that if you want to test your work using RARS, you can, but first try to convince yourself of the correctness of your work without actually writing a RARS program.

What to Include in Your PDF Submission

There is nothing to include. Solutions will be posted on the no later than February 28.

Exercise E: Integer subtraction examples

Read This First

An 8-bit adder can compute $a - b$ by inverting the bits of *b* and supplying a value of 1 as the carry in to the least significant bit. For example, if *a* is 0000_1001 (9 in

base ten) and b is 0000_0100 (4 in base ten), here's how the subtraction works:

```

carry in:          11110111
-----
bits of a:         00001001
inverted bits of b: 11111011
-----
result:            00000101

```

Note that the correct result (5 in base ten) is produced.

What to Do

Like Exercise B, this is a pencil and paper exercise. Don't use a computer or a calculator.

For each of the following pairs of hexadecimal patterns, write 8-bit binary patterns, and find the 8-bit subtraction result $a - b$. As in Exercise B, identify cases where signed overflow has occurred and cases where unsigned overflow has occurred.

PART	a	b	PART	a	b
I	0x20	0x8f	III	0xac	0xa5
II	0xc8	0x6e	IV	0x26	0x27

Show the work you did to generate the subtraction results and explain in each case how you decided whether signed overflow and/or unsigned overflow has occurred.

What to Include in Your PDF Submission

Include answers and supporting work for all the subtractions.

Exercise F: Review of D Flip-Flops

Read This First

D flip-flops (DFFs) were covered in ENEL 353. They are very important building blocks for the processor designs that we started studying in lectures just before Winter Break Week.

The critical thing to understand about a DFF is that its state Q can change *only* in response to an *active edge* of its *clock* input.

- For positive-edge-triggered DFFs (the kind we saw a lot in ENEL 353) the active edge of the clock is the *rising edge*—a low-to-high transition;
- for negative-edge-triggered DFFs (which we did *not* see much in ENEL 353), the active edge of the clock is the *falling edge*—a high-to-low transition.

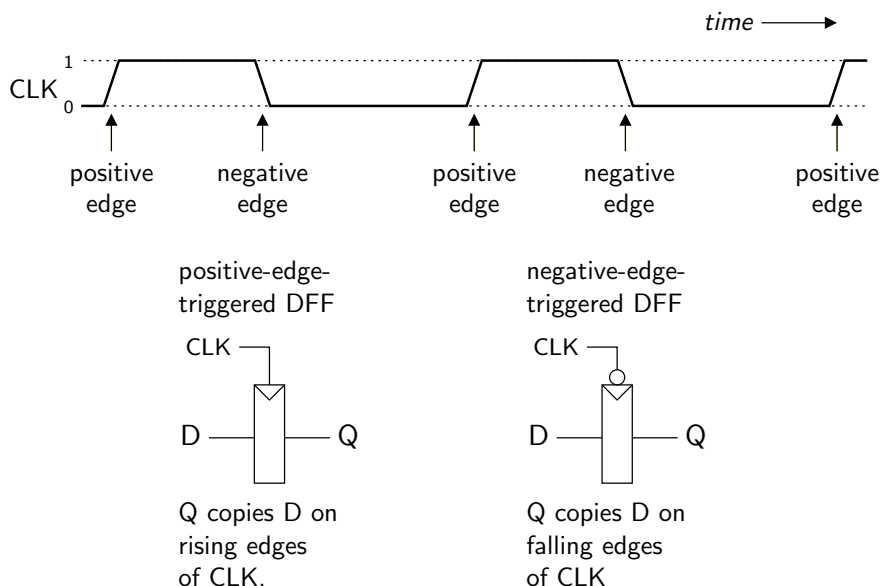
Behaviour of DFFs is summarized in Figure 2.

Knowing how a DFF responds as a “black box” to its clock and D inputs is essential in understanding sequential logic designs. On the other hand, knowing what is *inside* a DFF is not very important at all, unless your job is designing DFFs to meet specifications such as high speed, low chip area, and low power consumption.

Figure 3 shows a classic design for a DFF that was presented in ENEL 353. Don't worry if you don't understand how it works—that's not important in ENCM 369. (Note also that DFFs in modern CMOS integrated circuits don't use NAND gates at all—typically DFFs are built from inverters and devices called *transmission gates*, which we did not study in ENEL 353.)

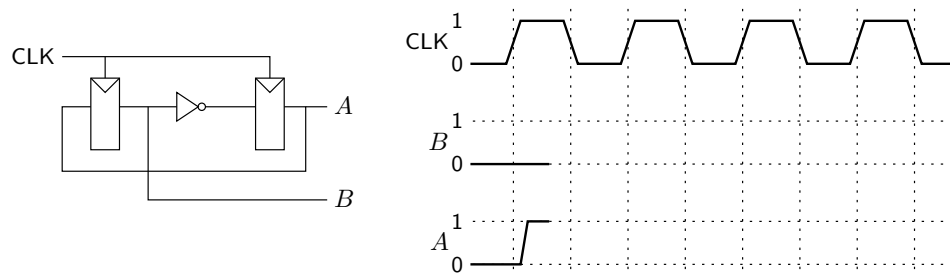
Figure 2: Summary of D flip-flop behaviour. *Nothing about digital logic design is more important than this!* Getting this wrong in any way should be considered to be just as serious a mistake as writing an incorrect truth table for an inverter, an OR gate, or an AND gate.

A clock signal ...



Very shortly *after* an active clock edge, Q takes on the value D had just *before* that clock edge; Q then *holds* that value until the next active clock edge.

Example problem. (This was a problem on the December 2012 Final Exam in ENEL 353.) Complete the timing diagram:



Solution to example problem. Before drawing waveforms, a very small amount of analysis is required:

- The DFFs are positive-edge-triggered, so the dotted lines dropping from negative edges of the clock are irrelevant and should be ignored.
- For the FF that has A as its output, the D input is \overline{B} . (I am assuming you have not forgotten what an inverter does!) That means that just after a positive clock edge, A will take on the value \overline{B} had just before that edge. In the notation of the textbook used for ENEL 353 and ENCM 369 used last fall and for the previous several years— $A' = \overline{B}$.

Figure 3: Positive-edge-triggered D flip-flop built from two NAND-based D latches. **Knowing exactly how this circuit works will NOT help you at all in ENCM 369!** It is *infinitely more important* to understand what is meant by the text in Figure 2!





Figure 4: Solution to example DFF problem. Moving left to right, each time you encounter a positive clock edge, use values of signals just *before* the edge to find the values the FF outputs will take on *after* the edge.



- For the FF that has B as its output, A is wired directly to the D input. So $B' = A$.

This analysis leads to the solution in Figure 4.

What to Do

Print a copy of page 9 of this document, or load it into software that lets you draw on the page.

In the spaces beneath each of the two given circuits, write next-state equations for the circuit. In Part I, that means equations giving Q'_1 and Q'_0 in terms of Q_1 and Q_0 . In Part II, that means equations giving Q'_0 , Q'_1 and Q'_2 in terms of Q_0 , Q_1 , Q_2 and the input A .

Then complete the given timing diagrams. Be very careful to check which clock edges are active for the given DFFs.

What to Include in Your PDF Submission

Include your completed worksheet.

